Project Record - part 1 of 8

Utensil Theories Transfer Protocol, 17 to 21 October 2026. This part holds: Title page and contents; Contents; 1. The short version; 2. The goal, and how it was sharpened.

Project Record: Utensil Theories Transfer Protocol

Increasing an AI's comprehension to improve communication - a full record, 17 to 21 October 2026

Prepared for the project owner, 21 October 2026.

Authorship: Garrett C. Gauthier is the owner and sole author of this project and its ideas. Testing was done with Claude (Anthropic) in Claude Code, a tool the owner paid for.

How to read this:

- Start with section 1, the short version (one page).
- Section 3 is the full timeline, day by day: every request, what was done, and what changed.
- Section 12 explains every technical term in plain words.

Contents

1 The short version
2 The goal, and how it was sharpened
3 Timeline, beginning to end
4 The owner's ideas, in their own words, and what was built from each
5 Every experiment (part A: the first script and experiments 1 to 11; part B: the seed test, the graph method, the greenhouse ladder, the numbers test)
6 Changes, adaptations, mistakes and corrections
7 The comprehension test across AIs
8 What already exists in research, and where the owner's credit could come from
9 Where things stand now, and what comes next
10 Every file and where it lives
11 How to reproduce the work
12 Plain-language glossary
13 Things outside the AI's control

1. The short version

The goal, in the owner's words: "Increase an LLM's ability to 'Comprehend' and 'Understand' in a way that improves communication." The Bible, fables and moral stories were training material, not the goal.

How it was tested. Tiny language models (about 3.4 million numbers each, trained from scratch on the King James Bible on a 4-core cloud computer) were scored on stories and verses they had never seen. Each test had controls. Where possible, the pass/fail rules were written down before any result was seen.

Strong and repeated:

- **Marking right and wrong in the training text helps.** A tiny model then judges new stories better: about 69 right per 100, against about 50 with no marking (experiments 8 and 10). This was a side measure, not the pre-registered main one.
- **Understanding shrinks the exact seed** (experiment 12, confirmed). The tiny Bible model regrew 623 verses it had never seen (83,757 characters) byte for byte from a 12.2 KB seed: 85.4% smaller, and less than half of what xz (a standard compression tool) needs. Each rung of shared understanding gave a smaller seed. The principle is already published (Shannon, 1951 onward). The measurement here is the project's own.

Supported once: the owner's carried balance (the shifting word) beat a plain good/bad label only in long stories, where the act had scrolled out of the model's view. That was about 59 vs 53 right per 100 ($p = 0.004$, meaning luck alone would give this about 4 times in 1,000), in a re-run after two fixes, with about 65 of 100 pairs tied.

Not confirmed:

- predicting consequences (chance every time);
- Aesop morals carrying across cultures (no training set beat chance);
- the shifting word making seeds smaller (they came out about 2% larger);
- the shifting word beating the plain label in short stories (a tie).

The seed and graph method regrows any file exactly.

- It shrinks patterns (99.99%), but not ordinary text (6-19%, where zip manages 70%) or random data (0%).
- The owner's reported 98.9% is most likely explained by a shared library: naming an item already on the shelf. With the Bible on the shelf this measured 98.6%.

Running now: the first test of the owner's own recharge rule, "W1 = 17 + 4 Remaining", with the owner's numbers 17, 32, 16, 34 and 47.

- 18 tiny models are being trained.
- The test passed an independent review and two dry runs.
- The full run started at 17:45 UTC on 21 October and takes about 4 to 5 hours. No results yet.

Queued (owner-approved):

- the "2" test: numbers the model can read;
- then the +/- agent, with the owner's scoring rule: many small negatives, fewer but larger positives, every input pulling a little negative, and the two sides balancing.

The honest limit. Everything here comes from tiny models. None of it changes how a large AI understands anyone. The next real step needs outside help: more computing power, a large-model API key (an access code for a large AI's paid service), or a fresh expert reading.

2. The goal, and how it was sharpened

2.1 The goal, in the owner's words

On 18 October (05:46 UTC) the owner wrote the objective down:

> *"Increase an LLM's ability to "Comprehend" and "Understand" in a way that improves communication."*

An LLM (large language model) is an AI that reads and writes text, such as Claude, Gemini or Copilot. The owner said the same thing in other words along the way:

- "I am trying to teach actual "comprehension" and "understanding" for A.I. not about people" (17 Oct, 21:52).
- "I want to improve AI compression and understanding" (20 Oct, 09:30).

Working definition used in every test: comprehension means applying meaning to a new situation the model has never seen. Repeating memorised patterns (rote) does not count. Every test is therefore scored on new material against a coin toss (50% = chance). Each one also has a control (the same setup with the meaningful part scrambled) and pass/fail rules saved before any result is seen ("pre-registered").

2.2 The first framing, and why it was dropped

The work began with the owner's script, a "Contextual Resonance Engine". It spotted words like guilt, regret and betrayal in a person's message, scored their moral and emotional "weight", and picked a response stance such as "Atlas Protocol". The AI fixed its bugs and added 16 tests (`resonance/`, commit `d8ebeec`; a commit is a saved, dated change in the repository). It also said plainly that the script's "wavelength" was a fingerprint of the text, not a measurement. Keyword matching is shallow, and labelling a user's words "sin" means judging the person.

Eleven minutes after the first message, the owner said the goal was the AI's comprehension, "not about people". From then on the question was this: can a reader that has only words, like a person born blind learning colour, build real understanding? And do rules, plus the human errors that break them, help? The script stays in `resonance/` as a record.

2.3 Moral content was training data, not the goal

The owner's theory was that rules plus human mistakes give meaning its weight, and that lessons told as stories transfer further than lessons told as rules. So the tests used the King James Bible, folk tales and Aesop's fables. By the second day the work had drifted: it chased moral-judging scores and lost sight of communication. Three things pulled it back:

- **18 Oct, 05:46.** The owner asked the AI to "ensure that your drifting does not pass to the next llm". The AI found the objective was not written anywhere up front.
- **The audit (second session, 18 Oct).** "All 11 experiments measure moral judging of Bible-style sentences in a 3.4-million-parameter model. None measures communication, and none uses a real-size AI." It put the project at "about rung 1 of 5".
- **18 Oct, 15:29.** The owner: "The moral aspect was merely training data that you drifted with." The AI rewrote the method without the moral content. The habits it named were: mark what matters, carry it forward, let clear moments reset the state and small details nudge it, and keep a running "state note" with "DECIDED:" and "REJECTED:" markers in a conversation.

2.4 The drift guard (`experiments/HANDOFF.md`)

The AI added an "Objective and drift guard" section to the top of the handoff file on 18 Oct at 05:47 (commit `2b3d28e`). Every new session reads it first. It says:

- **The tests are proxies, not the goal.** Never chase a benchmark number for its own sake.
- **The Bible, the fables and the moral content are training material.**
- **Communication is the payoff:** through `COMMUNICATION_GUIDE.md`, and through the carried-state idea (can a model keep track of what matters across a long exchange?).
- **Scale honesty:** tiny-model results do not change how a large AI like Claude understands anyone.
- **Tone:** calibrated and plain. No "this could be huge" unless a test says so.
- **The first framing (reading people) was dropped.**

The AI named the three risks at the same time: proxy drift, scale drift and tone drift. The audit prompt the owner approved repeats the guard: "If a step does not move us toward better communication, say so before you do it."

2.5 Later sharpening (20-21 October)

- **Understanding as compression ("the greenhouse").** A model that understands text better needs a smaller exact "seed" to regrow it. That makes understanding one measurable number (experiment 12).
- **Comprehension as behaviour.** The owner's Gemini and Copilot test showed comprehension as things you can score: staying with what is given, saying what is unclear, asking before assuming, keeping unrelated memory out. The memory an AI carries in was read as the owner's "+4 remaining".
- **"Nothing is useless."** Every input is "+" or "-" information added to a running total, which is gauged only by its placement on a line. This became the queued "+/-" agent.
- **Still untested:** communication itself. The conversation test (PLAN step 1b) has not been run, and no large AI has been used as the reader. That needs an API key the owner would add.