Project Record - part 2 of 8

Utensil Theories Transfer Protocol, 17 to 21 October 2026. This part holds: 3. Timeline, beginning to end.

3. Timeline, beginning to end

How to read this timeline:

- Times are UTC. The owner's local date may differ.
- "#n" is the owner's message number (59 in all).
- Short codes such as `d8ebeec` are commits: saved, dated changes in the repository.

Tags:

- **[DIRECTION]** the request changed direction;
- **[CORRECTION]** the owner corrected the AI;
- **[PAUSE]** work stopped on purpose;
- **[LIMIT]** a usage-limit stop;
- **[RESTART]** a step was re-run or resumed;
- **[BLOCKED]** a safety check stopped the reply.

Earlier commits (`b8057a8`, `75bbff6`, `690d671`, July-August) belong to the owner's separate zine project, which is out of scope.

17 October 2026 (Saturday)

- **21:41 #1** - The owner asked to "Evaluate and enhance" the resonance script. The AI ran it and confirmed the bugs: missed word forms, ignored negation, one word triggered "grace", one stance could never fire, and the "wavelength" changed every run. It rebuilt the script with 16 tests (`d8ebeec`).
- **21:47 #2** - "Is this feasible, am I onto something?" The AI said yes, as a tone guard, but it was half tool and half metaphor.
- **21:50 #3, 21:51 #4** - The owner pasted a chat with another AI about "the architecture of human error" (Commandments, unforgivable sins, "A mistake ... is the calibration point") and asked the AI to separate the idea from the task. The chat held personal reflections, which are not reproduced here. The AI warned of the other AI's "agreement loop" and mapped the idea as rule -> deviation -> weight -> direction -> grace.
- **21:52 #5 [DIRECTION]** - "...for A.I. not about people".
- **22:02 #6 [DIRECTION]** - The blind-colour question; the Bible as "a road map to understanding"; "have faith and give it a try. If it doesn't work we will try something else." Experiments 1-4 on the King James Bible found:
  - colour matched 10 of 14 (chance 17%);
  - rule verses alone were no better than other verses;
  - "covet" weakened by about half without the rules.

Commit `ec22991`.

- **22:12 #7** - "try the consequences test". Experiment 5: chance (45-57%) on 20 new story pairs (`939d60f`).
- **22:19 #8** - "try the tiny transformer". Experiment 6: a tiny modern-design model (a "transformer") learned Bible language in 12 minutes, but scored chance on the stories (45-55%) (`df94814`, `1a4dfac`).
- **23:00 #9 [DIRECTION]** - The owner's shifting word: "[w1=17 -<next word ... (-13)> = +4 Remaining]". It resets to +17 or -17 on a clear right or wrong, its spellings vary like "to, too, two, 2", and it should "shift the coin toss". The owner also noted their numbers 17, 32, 16, 34, 47. The AI built "velanok": five syllable-rotation spellings and a fading balance. It switched off the owner's re-adding of 17 ("17 + 4"), because the balance would grow forever. Experiment 7: 68% vs 52% for the scrambled control, on 20 pairs (`f676eed` to `c918cb7`). This was later shown to be luck, and no rules had been saved first.

18 October 2026 (Sunday)

- **00:46 #10** - "Let's confirm the result". The owner also asked what it would mean for them and whether it could earn anything. The AI saved the rules and 100 new pairs before training (`5538b46`). Its opinion: a small, real finding if confirmed; its value is reputation rather than direct earnings; credit the idea to the owner and the code to the AI.
- **01:41 #11 [DIRECTION]** - "the stories of the knowledge keepers and the myths of old"; "the theory of relatively within wording". The AI built a library of 12 public-domain collections from 9 cultures, about 727,000 words (`7c4a144`).
- **01:52 #12** - Run the Aesop test, then hand off to a new instance so "we dont lose progress". "This could be huge actually". The AI noted that tiny models do not change how it understands the owner. It saved the Aesop rules (`87cccf9`), which reviewers fixed twice before training (`194fc7c`, `910e9cb`). It also wrote the communication guide, handoff, prompt and plan (`b58b583`, `cde3dac`, `6143fa8`).
- **01:58 #13** - "keep going and tell me the verdict".
- **02:02 #14** - "When will I run out of usage?" The AI cannot see limits; everything was saved.
- **02:12 #15** - How does this compare with today's best? Experiment 8 was **not confirmed** for consequences (51/100, p = 0.46). Judging was strong as a side measure: 67, 72 and 67 right out of 100 (majority vote 69/100, p = 0.0001). It is far from today's large AIs.
- **02:14 #16 [CORRECTION]** - "evaluate it as methods". The AI agreed and saved rules for a head-to-head of the shifting word, a plain good/bad label and a 3-spelling carry (`6e365ef`). The audit later found that judging was chosen as the main measure after a peek.
- **02:21 #17** - "even if its .5% ... how much". About 15 more right per 100, not yet confirmed.
- **02:35-04:16 (no new requests).**
  - Experiment 8 was written up (`a5f1a6b`).
  - Experiment 10 was a tie: shifting word 68.7%, label 67.7% (`9c6fb86`).
  - Experiment 11 (long stories) had its rules saved at 03:44. It was then changed twice mid-run **[RESTART]** (`c93f383`, `6fc220d`).
  - It was written up as "confirmed": 59 vs 53 per 100 (`01b9050`). That was later corrected to "supported once".
  - The word "carries perfectly but updates poorly" (`11982b6`).
- **05:46 #18 [DIRECTION]** - "launch the new session now", with the drift instruction. The AI added the drift guard (`2b3d28e`) and launched a second session.
- **05:48-06:53, second session** (the owner's messages there are not in this record) - It audited the work: the numbers matched the raw files, but some wording was too strong (`84e2ba4`). It wrote the owner's file rules into `CLAUDE.md` (`3360d93`), corrected "private" to "public repository" (`a67ebdf`), and sent this session an owner-approved audit prompt (`685daf7`).
- **06:11 [LIMIT warning] [PAUSE]** - A 7-day usage warning appeared (it resets 24 Oct). This session went quiet while the Aesop models trained.
- **07:33-07:39** - Aesop: all four checks **not confirmed**. Folk tales 52.3%, novels 54.1%, shuffled 53.1% (`7fdad39`, `29a94aa`). The audit corrections were applied (`2a58b61`, `6bfa507`). The results folder was delivered and `CLAUDE.md` was copied to this branch (`d93b6f7`, `49e10f0`). The audit branch merged the results (`3ee90de`, `b806b95`). **[PAUSE]**: the conversation test waits for the owner to approve its rules.
- **15:29 #19 [CORRECTION]** - "The moral aspect was merely training data that you drifted with". The owner asked for copyable plain text. The AI wrote a 6-part plain guide in the chat.
- **16:02 #20, 16:04 #21** - The owner pasted another AI's "Data Filtering & Methodology Report" (v1.0, then v2.0). The AI found that v1.0 was a summary of its own write-ups: it overstated results, mixed up two experiments and drifted from communication. v2.0 was better, but three statements were still wrong. The AI gave correction lines to paste back.

19 October 2026 (Monday)

- No owner messages and no commits.

20 October 2026 (Tuesday)

- **08:13 #22 [DIRECTION]** - The owner uploaded 11 files (other AIs' reports, a compression design called EBEC, notebooks, a prototype, and profile documents with personal details) and asked for "a full technical report in .doc format" with all code.
- **08:17 #23** - The owner stopped a tool call: the AI was starting to run tests.
- **08:35 #24 [CORRECTION]** - "I didn't say run tests? Did I?" The owner wanted a report detailed enough to regrow everything in a fresh chat: a "seed" plus "grow instructions", like algebra's "show your work". "Do not fail to generate my doc". The AI delivered `TECHNICAL_REPORT_2026-10-20/`: a real .doc and .docx, the seed, 10 grow steps, a source evaluation and all the code (`8c08f41`). Personal details were left out.
- **08:45 #25** - Could data seeds save electricity and add security? The AI said compression is real. But new information cannot be shrunk, regrowing costs power, and a "dead seed" is security through obscurity. The fix: compress, then encrypt.
- **08:56 #26 [CORRECTION]** - "You dont understand my friend how it actually works". The owner reported 100% reconstruction and 98.9% reduction with Copilot and Gemini, with no LLM needed. The AI built a fair, AI-free test, `seed_test.py` (`283b4b6`). With zip, random bytes shrink 0%.
- **08:57 #27** - "Use a graph system lol". Graphs suit patterns, not random data.
- **08:59 #28, 09:00 #29, 09:03 #30** - "add the graph method". The owner interrupted to explain: a fixed "page of binary" built from solving an expression, plus library factors "strawberry and orange verse cucumber".
- **09:06 #31** - "Run it. Your close". The AI built three readings, each with an independent checker (`6cd7332`).
- **09:10 #32, #33** - "tell me the verdict".
- **09:14 #34** - "Are you surprised?" No. The same idea lives in game world seeds and in Kolmogorov complexity (the shortest description of data).
- **09:20 #35** - A palm seed grown in "Antarctica": different operating systems. The AI said exact rules work everywhere if they use whole numbers and raw bytes. The graph-method drafts so far were saved (`b6397c1`).
- **09:30 #36 [DIRECTION]** - The greenhouse story: soil, humidity, sun, then a seed. "I want to improve AI compression and understanding", and credit. The AI launched the greenhouse ladder, a dig through the chat and a fact check. Graph-method drafts were saved (`8f705c6`).
- **09:45 #37** - "map out these points as if expression on a face". The AI posted a progress map with faces. A graph-method update was saved (`93890e4`).
- **09:55-10:13 [LIMIT]** - Session limit (reset 13:10 UTC).
- **16:13 [RESTART]** - Graph verdict (`d7b1323`). Every file regrew exactly, but only patterns shrank: pattern 99.99%, Bible text 6-19%, random 0%.
- **17:08** - Experiment 12 (`8dff9ba`). The tiny Bible model regrew 623 unseen verses from 12.2 KB instead of 83.8 KB (85.4% smaller, exact). The shifting word did not help.
- **19:09 #38 [DIRECTION]** - "Add them and test the nunbers "17,32,16,34"=?0=-0=+0=00=["47,"="17+4"]". The AI saved the 15 trained models to GitHub (`33a893d`) and started a design team.
- **19:19 #39 [CORRECTION] [RESTART]** - "the +4 was a way to direct you to a way to try testing the numbers where I feel like they belong." The design restarted with the numbers in 17's place in "W1 = 17 + 4 Remaining", in the owner's order.
- **19:24 #40** - "Which 2,to,too or two feels like a missing clue?" The AI picked "2": the model's reader drops digits.
- **19:26 #41** - "add the 2 test after this one". Queued (`7695b97`).
- **19:34 #42 [DIRECTION], 19:37 #43** - An astrology/numerology formula, as a test of how other AIs comprehend. The AI wrote a kit: the formula, a shuffled control, a follow-up question and written predictions.
- **19:41 #44, 19:43 #45** - The owner brought back the answers. Gemini (memory on) gave the shuffled control a confident meaning and pulled in the owner's other chats. Copilot stayed closer to the text, but still read meaning into the control. The answers were kept out of the public repository because they held personal context.
- **19:47 #46 [CORRECTION]** - "We need to understand the "17+4" they used". The AI read each answer as the prompt (17) plus carried memory (+4).
- **19:52 #47** - Gemini with memory off stayed with the words. The +4 was real and, this time, it hurt.
- **19:57 #48** - How would this AI answer? It gave its own answer (with a large +4, since it wrote the prompt) and predictions for a fresh copy of itself.

- **20:02 #49 [CORRECTION]** - "Nothing is useless." "You Do Not Know the Answer". "a sharp left away from logic?" The AI agreed and proposed a random-numbers control. That control was not adopted, and the AI later called its framing wrong. The draft script was saved (`12f2051`).
- **20:02 #50 [PAUSE]** - "Pause tasks". Stopped before any training.
- **20:13 #51 [CORRECTION] [DIRECTION]** - "Everything adds and subtract from understanding and comprehension." The owner shared personal context for why every input is + or - information and why recharging in a chosen direction matters (not reproduced here). Their instruction was an agent that sorts each input into + or -: "dont do this step yourself but make sure it happens", and never know the true number, only its "placement on that line". The AI said "You're right, and I was wrong" and queued the agent.
- **20:22-21:06 [LIMIT]** - Session limit (reset 21:10 UTC).

21 October 2026 (Wednesday)

- **01:35** - The AI confirmed everything was paused and saved.
- **01:41 #52 to 01:53 #56 [BLOCKED x5]** - The owner asked for a plain review from beginning to end:
  - reasons "for", "against", and anything overlooked;
  - every source and every "+4" the AI noticed in its replies;
  - a word-for-word rerun and comparison (#53);
  - "why a review ... is a risk" (#54, a correction);
  - a "safe prompt" and the usage back (#55; the AI has no access to billing);
  - a full word-for-word review (#56).

The owner also noted that the AI is listed as co-author on commits. All five replies were stopped by an automatic check ("safeguards flagged this message", detail "[reasoning_extraction]"). The AI could not see why. One unconfirmed reading is that the wording looked like a request for its hidden reasoning. No work was lost.

- **01:55 #57** - "a complete understanding as if I was a manager and you are the worker". This one got through (01:59). It held every request and what was delivered, the results by strength, the untested items, the AI's own errors (its "+4s"), and for/against/overlooked.
- **02:05 #58 [RESTART]** - "Continue on with testing". The numbers-test dry run passed (02:06-02:08), and the folder `NUMBERS_TEST_2026-10-21/` was saved (`0e4574a`, `5d20314`). A 4-reviewer check is running. **No results yet.**
- **02:15 #59** - This report.

Totals:

- 59 owner messages;
- 2 usage-limit stops, plus one 7-day warning;
- 3 deliberate pauses;
- 2 owner interruptions;
- 5 safety blocks.

21 October 2026, later (added after this timeline was drafted)

- **02:21 #60** - The owner suggested emailing some people a link and inviting them to the repository.
  - The AI explained that the repository is already public, so a link is enough. Collaborator invites (which let people edit) are made by the owner in GitHub settings.
  - It offered to write a Gmail draft for the owner to send once recipients are named. None have been named yet.
- **02:22 #61** - The owner observed that the branch's diff (its list of changes against the main version) has stayed large and similar while progress has slowed: "We need other people, someone else to know."
  - The AI measured the branch: +156,235 lines against `main`, 141,791 of them on 20 October, mostly copies of data files (the Bible text and one results file, each stored twice). Line counts are not progress.
  - It agreed the real limit is scale (tiny models on 4 CPU cores). It proposed outside help: computing power, a large-model API key and a fresh critical reading. It offered a one-page summary for newcomers and an outreach draft.

- **Early 21 October** - The session hit its usage limit (reset 06:30 UTC). It stopped this report and the numbers-test review part-way.
- **13:07 #62** - The owner shared another AI's Python blueprint, "ComprehensionSeed": a living memory of concept weights that are added to and subtracted from, with human errors treated as signal and a look beneath the literal words.
  - The AI's reading: it is an outline, and the seven functions that would do the hard part are undefined. It matches the owner's thesis and the planned +/- agent.
  - It has three flaws the experiments already showed: weights never fade; contradictions are only checked for new concepts; nothing is ever removed.
  - One framework name in it contains a person's full name, which is kept out of this public record; it is called here "the owner's error-encoding method".
- **13:12 #63** - The owner gave the scoring rule for the +/- agent:
  - **Positive:** moral high ground, telling the truth, highly reasonable information.
  - **Negative:** questionable morals, nonsense, misspelling.
  - "Should always pull a negative." More negatives than positives, but the positives are larger, and it balances.
  - It is recorded in `HANDOFF.md` (`a8a0cd7`). The AI measured today's word values as symmetric (36.0% vs 37.4%; average sizes 2.31 vs 2.27).
  - The size of a positive against a negative is still the owner's to choose.
- **13:10 to 13:30** - The review and this report resumed, then hit the usage limit again (reset 17:00 UTC).
- **17:37 to 17:45** - The AI did the remaining run-safety checks itself and applied 10 small fixes that leave the frozen design unchanged. A second dry run passed. **The full 18-model run started at 17:45 UTC** (`55f6817`).

Updated totals: 63 owner messages; 4 usage-limit stops (2 on 20 October, 2 on 21 October).